

GPU programming made easy with OpenMP

Aditya Nitsure (anitsure@in.ibm.com)

Pidad D'Souza (pidsouza@in.ibm.com)



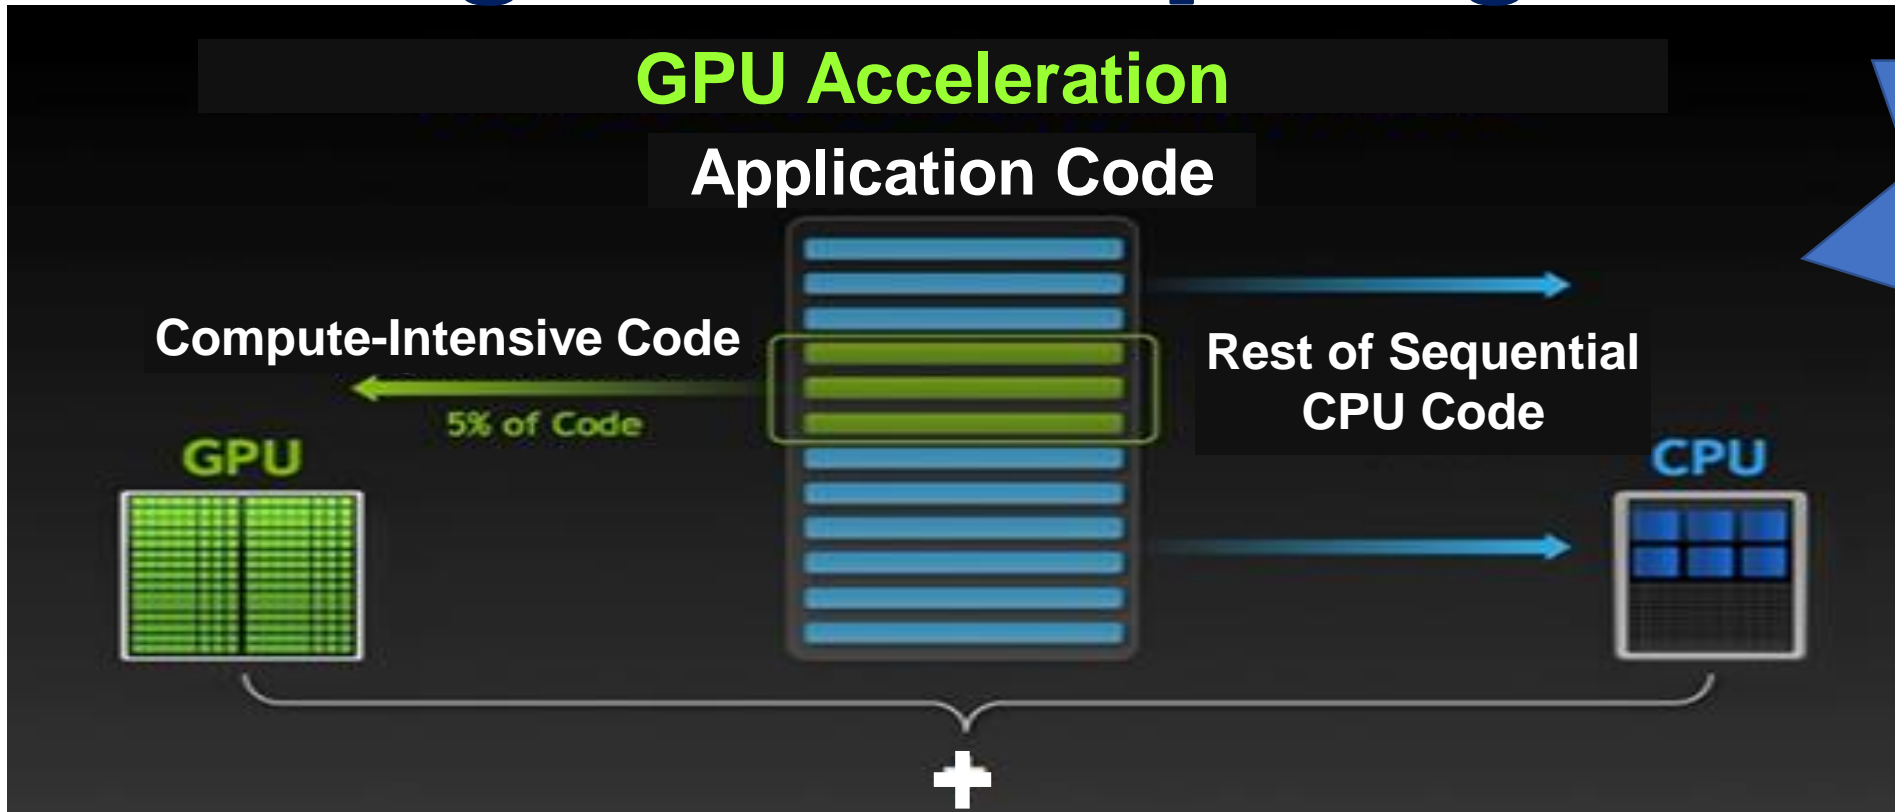
Tutorial Overview

- Heterogeneous system overview
- Introduction to parallel programming
- OpenMP programming on CPU and GPU
- Profiling and monitoring
- HPC application performance

- Hands-on

Heterogenous Systems

Heterogenous Computing





Maximize
performance
and energy
efficiency

CPU

- Large and broad instruction set to perform complex operations

GPU

- High throughput – Massive parallelization through large number of cores
- Specialized for SIMD/SIMT




Heterogenous System configurations

Todo

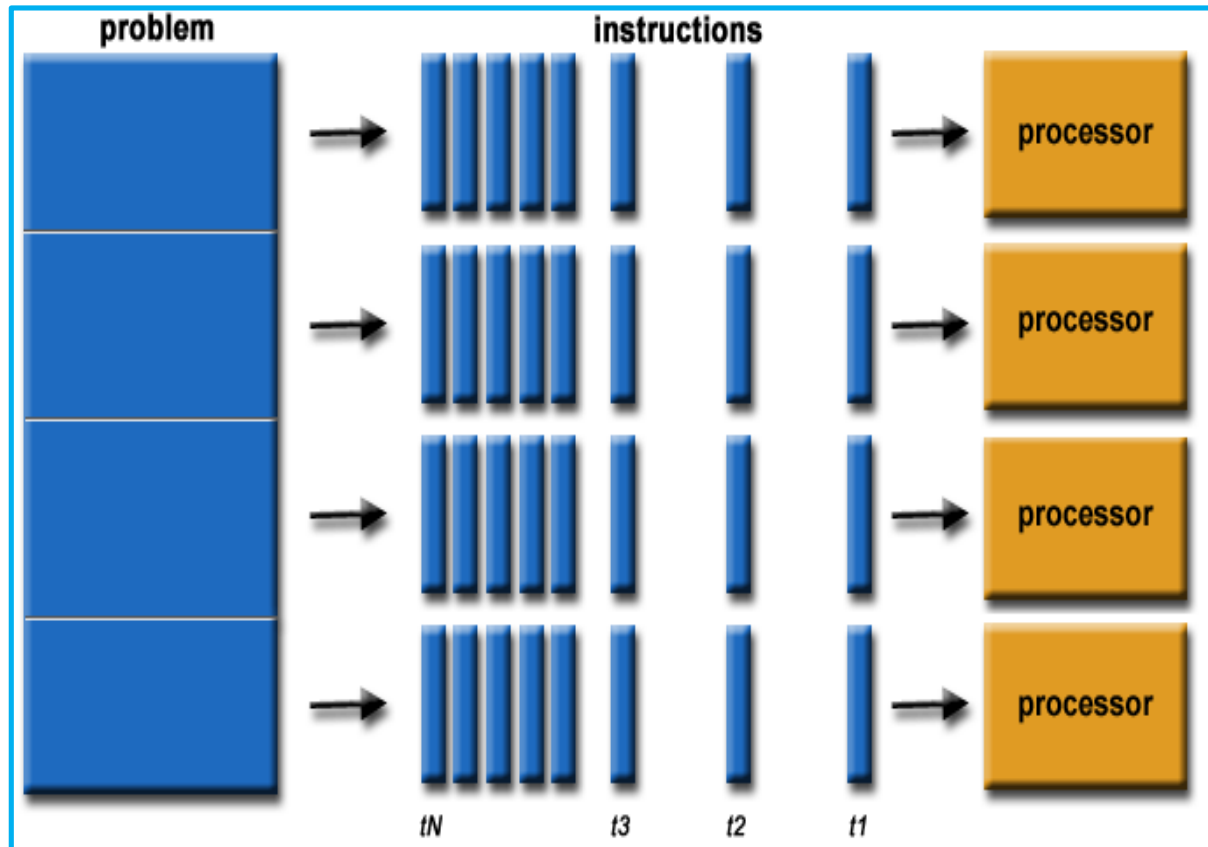
Introduction to Parallel computing

What is Parallel Computing?



Sequential Computing

- Sequential execution of **series of instructions**
- Only one computing resource







Parallel Computing

Simultaneous use of **multiple compute resources to solve a computational problem**

- Divided problem into discrete sub-problems
- Execute sub-problems in in simultaneously on different compute resource
- More than one compute resource

More slides to follow

Memory Architecture



Proprietary and Open Source Compiler Offerings supporting Acceleration Enabled Programming Models



Key Features:

- Gives direct access to the GPU instruction set
- Supports C, C++ and Fortran
- Generally achieves best leverage of GPUs for best application performance

OpenCL

MPI – Message passing



Key Features:

- OpenMP 4.0 introduces offloading and support for heterogeneous CPU/GPU
- Leverage existing OpenMP high level directives support




Key Features:

- Designed to simplify Programming of heterogeneous CPU/GPU systems
- Directive based parallelization for accelerator device

Applications of Parallel Computing



Science &
Engineering
Artificial Intelligence



Analytics
Realtime simulations
Modelling



OpenMP

Monitoring and Profiling tools

Monitoring and Profiling tools

Monitoring

- mpstat, vmstat – CPU and memory utilization
- numastat – numa memory statistics
- top/htop – real-time view of system usage

Profiling

- Perf record/report – CPU profiling
- nvprof – GPU profiling

[aditya@hpc] **numastat** -m | grep "Node\|MemUsed"

	Node 0	Node 8	Node 252	Node 253	Node 254	Node 255	Total
MemUsed	104872.31	57992.06	11.19	11.19	11.19	1455.19	164353.12

CPU memory

GPU memory

Monitoring and Profiling tools

nvidia-smi

```
[aditya@hpcnw4 ompeval]$ nvidia-smi
Sun Apr 21 03:03:12 2019
```

NVIDIA-SMI 396.64				Driver Version: 396.64			
GPU Name Persistence-M				Bus-Id Disp.A			
Fan Temp Perf Pwr:Usage/Cap				Memory-Usage GPU-Util Compute M.			
0	Tesla V100-SXM2	On		00000004:04:00.0	Off	1539MiB / 15360MiB	100%
N/A	42C	P0	153W / 300W				Default
1	Tesla V100-SXM2	On		00000004:05:00.0	Off	11MiB / 15360MiB	0%
N/A	38C	P0	37W / 300W				Default
2	Tesla V100-SXM2	On		00000035:03:00.0	Off	11MiB / 15360MiB	0%
N/A	35C	P0	36W / 300W				Default
3	Tesla V100-SXM2	On		00000035:04:00.0	Off	11MiB / 15360MiB	0%
N/A	41C	P0	38W / 300W				Default
Processes:				GPU Memory Usage			
GPU	PID	Type	Process name				
0	54063	C	./matmul_gpuoffload_cl	1528MiB			

Also check “nvidia-smi –query-gpu” more monitoring options

Monitoring and Profiling tools

nvprof


- The *nvprof* is command-line profiling tool which enables you to collect and view profiling data
- Todo

NVVP (NVIDIA Visual Profiler)

- The Visual Profiler displays a timeline of your application's activity on both the CPU and GPU so that one can identify opportunities for performance improvement.
- todo

HPC Application performance with OpenMP GPU offloading

LULESH : OpenMP 4.5 GPU Offload



LULESH

A shock hydrodynamics mini-app for computer simulations leveraging high performance computing of a wide variety of science and engineering problems that describes the motion of materials relative to each other when subject to forces.

Chart to follow

References

- GPU programming made easy with OpenMP on IBM POWER
(<https://developer.ibm.com/articles/gpu-programming-with-openmp>)
- Offloading computations to the NVIDIA GPUs
(https://www.ibm.com/support/knowledgecenter/en/SSXVZZ_16.1.0/com.ibm.xlcpp161.linux.doc/proguide/offloading.html)
- Parallel Computing(https://computing.llnl.gov/tutorials/parallel_comp)

Thank You